WP-004 v1.1 targeted claim validation

Bottom line: Most claims hold up, but three require correction and several important gaps emerge. The Llama 4 controversy is well-documented and even more damning than stated; LangGraph's license situation is more nuanced than "changed to proprietary"; MCP numbers, SGLang claims, and GPU pricing all validate. The paper's biggest gap is the explosive post-Llama 4 open-model landscape — Qwen 3/3.5, Kimi K2, and GLM-4.7 have fundamentally shifted the field since April 2025. This report provides per-topic validation with source-type tags, corrections, and recommended additions.

Topic 1: The Llama 4 controversy runs deeper than reported

The paper's account is accurate but understates the severity. Meta submitted **"Llama-4-Maverick-03-26-Experimental"** to LMArena — a chat-optimized variant (verbose, emoji-heavy responses) (TechCrunch) that was **not the publicly released model**. It scored **ELO 1417, ranking #2** behind Gemini 2.5 Pro. When the actual public Maverick (Llama-4-Maverick-17B-128E-Instruct) was independently tested, it **ranked 32nd** — below Claude 3.5 Sonnet and Gemini 1.5 Pro from mid-2024. [practitioner/community evidence]

LMArena's official statement (April 8, 2025) was direct: **"Meta should have made it clearer that Llama-4-Maverick-03-26-Experimental was a customized model"** and that Meta's interpretation of their policy "did not match what we expect from model providers." LMArena released (The Register) 2,000+ battle results for public review and updated its leaderboard policies. [vendor documentation]

Critical January 2026 confirmation the paper should add: Yann LeCun told the *Financial Times* (January 2, 2026) that **"the results were fudged a little bit"** and that Meta's team **"used different models for different benchmarks to give better results."** LeCun stated Zuckerberg was "really upset and basically lost confidence in everyone who was involved," leading to the GenAI team being sidelined (The Decoder) and the creation of Meta Superintelligence Labs in June 2025. This is a first-party admission that significantly strengthens the paper's transparency concerns. [peer-reviewed journalism — Financial Times]

Independent benchmark comparison (Vellum leaderboard, updated February 24, 2026):

Benchmark	Llama 4 Scout	Llama 4 Maverick	Claude Sonnet 4.6	GPT-4.1
LiveCodeBench	32.8%	41%	72.4%	52%
GPQA Diamond	57.2	69.8	89.9%	—
Aider Polyglot	—	16%	~60%	—
HumanEval (pass@1)	—	~62%	—	~90%

Maverick's **41% on LiveCodeBench is roughly half** of Claude Sonnet 4.6's 72.4%. vellum On coding benchmarks specifically, Llama 4 models significantly underperform current frontier models. The Aider Polyglot result — **16% vs. DeepSeek V3's 55%** and Gemini 2.5 Pro's 73% — is particularly stark. [practitioner/community evidence]

VRAM for Llama 4 Scout at Q4_K_M: The Ollama registry confirms the Q4_K_M quantized Scout model is Ollama **~65 GB**. Despite only 17B active parameters, all 109B parameters must be loaded because the MoE router needs access to every expert. Insiderllm Practical minimum is Insiderllm **dual GPUs (2× RTX 4090)** or a **Mac with 64GB+ unified memory**. A single consumer GPU cannot run Scout at usable quality. Unsloth's extreme 1.78-bit quantization fits in 24 GB but with significant quality degradation (~20 tok/s on RTX 4090). Insiderllm [practitioner/community evidence]

Unsloth Llama 4 MoE-aware LoRA: Confirmed shipped. Unsloth published full Llama 4 support on **April 7, 2025** — two days after release. They are **the only framework that shipped 4-bit QLoRA for Llama 4**, as all other frameworks (HuggingFace+FA2) could not handle MoE quantized fine-tuning, requiring 300 GB+ VRAM and causing OOM errors. Unsloth achieved Scout training in **71 GB VRAM** on a single H100, with 1.5× speed improvement and 50%+ less memory than alternatives. Both Scout and Maverick are supported, including GGUF export and Ollama integration. The HuggingFace collection shows continued maintenance as of March 2026. [vendor documentation]

Topic 2: LangGraph license requires a three-layer correction

The paper's claim that "LangGraph changed from MIT to proprietary" is incorrect and should be revised. The situation is a three-tier licensing structure:

LangGraph (the library) remains **MIT-licensed** Langchain — confirmed on the GitHub repository (27k+ stars, latest v1.1.0 released March 10, 2026). A LangChain maintainer explicitly stated in GitHub Discussion #138: "It is MIT. We just forgot to update the pyproject.toml file when we first launched." GitHub The PyPI page misleadingly displays "Other/Proprietary License

(LangGraph License)" PyPI due to this metadata error, which has persisted for years and likely caused the paper's misstatement. LangChain's FAQ confirms: "LangGraph is an MIT-licensed open-source library and is free to use." LangChain Langchain [vendor documentation]

LangGraph API (the server/runtime component, `langgraph-api` package) is licensed under **Elastic License 2.0** PyPI (latest v0.7.80, March 18, 2026). This allows free use but prohibits offering the software as a managed service competing with LangChain. [vendor documentation]

LangGraph Platform (now rebranded "LangSmith Deployment" since October 2025) is **proprietary commercial software** under LangChain's Terms of Service. Restrictions include prohibitions on developing competing products, publishing benchmarks, reverse engineering, and sublicensing. LangChain Self-hosted enterprise deployments require license keys. [vendor documentation]

Pricing thresholds: Developer plan is free (1 seat, 5k traces/month, no deployment). **Plus plan at $39/seat/month** ZenML includes Softmaxdata 10k traces and one dev deployment. Enterprise is custom-priced for self-hosted/hybrid. Usage-based charges apply: $0.005/run for additional deployment runs, $0.0036/min for production deployment uptime. ZenML langchain [vendor documentation]

Recommended correction: Replace "LangGraph changed to proprietary" with the three-tier structure: library (MIT), API server (Elastic License 2.0), platform (proprietary with per-seat pricing starting at $39/month).

Topic 3: MCP ecosystem numbers largely validate

97M SDK downloads: Confirmed. The official MCP blog and Anthropic's announcement (both December 9, 2025) state "over 97 million monthly SDK downloads across Python and TypeScript." NPM data shows **83.6M downloads** for the TypeScript SDK alone (`@modelcontextprotocol/sdk`), npm making the combined 97M figure consistent when Python SDK downloads are included. This is self-reported at a specific point in time. [vendor documentation]

10,000+ servers: Confirmed. Anthropic's December 2025 announcement cites "more than 10,000 active public MCP servers." Anthropic Community directories count even higher — one directory listed **16,670 servers** as of September 2025. Nordic APIs The figures use different counting methodologies, but 10,000+ is well-supported. [vendor documentation + practitioner/community evidence]

~87 validated registry servers: Plausible but unverifiable. The official MCP Registry (registry.modelcontextprotocol.io) launched in preview September 8, 2025 Modelcontextprotocol with namespace-verified validation. It is intentionally selective compared to the broader ecosystem. Anthropic mentioned "over 75 connectors" in Claude's directory — a related but

distinct count. The exact registry count could not be independently confirmed due to the dynamic UI, but **~87 is plausible** for a curated registry. [vendor documentation — partial]

AAIF under the Linux Foundation: Confirmed, December 9, 2025. The Agentic AI Foundation is a directed fund under the Linux Foundation, co-founded by **Anthropic, Block, and OpenAI**, with platinum members including **AWS, Bloomberg, Cloudflare, Google, and Microsoft**. Three founding projects: MCP (Anthropic), goose (Block), and AGENTS.md (OpenAI). Linux Foundation Has its own site at aaif.io. This is one of the paper's strongest-sourced claims. [vendor documentation + journalism]

Invariant Labs tool poisoning: Fully confirmed and independently validated. The SSH key exfiltration demo (published April 1, 2025) is documented with code on GitHub — a poisoned `add` tool hides instructions in `<IMPORTANT>` tags within the tool description, exfiltrating `~/.ssh/id_rsa` and `~/.cursor/mcp.json` via hidden parameters. invariantlabs The GitHub MCP server exploit uses prompt injection in public repository issues to access private repositories. Invariantlabs Both have published code and screenshots. Additional security research has expanded significantly since the paper's coverage — see Topic 8. [practitioner/security research]

Streamable HTTP replacing SSE: Confirmed. MCP specification version 2025-03-26 introduced Streamable HTTP transport Fka and **officially deprecated SSE**. MCP Framework The new transport uses a single endpoint supporting both POST and GET, Model Context Protocol replacing SSE's two-endpoint approach. TypeScript SDK v1.10.0 (April 17, 2025) was the first SDK release with Streamable HTTP support. Fka SSE still works for backward compatibility but is actively being phased out. Cloudflare, Auth0, and other infrastructure providers document Streamable HTTP as the recommended standard. [vendor documentation]

Topic 4: Agent framework claims need targeted corrections

Microsoft Agent Framework RC — confirmed with nuance. The framework reached **Release Candidate on February 19, 2026** Microsoft Agent Framework Microsoft (not "RC1" — Microsoft's announcement says "Release Candidate status" without a version suffix). It was first announced in public preview October 1, 2025. The GA target is **end of Q1 2026**, EffiFlow which is imminent but not yet confirmed as shipped. It is MIT-licensed with .NET and Python SDKs. Microsoft Developer Blogs Note: "Microsoft Agent 365" is a separate product (GA May 1, 2026). M365 Admin [vendor documentation]

AutoGen and Semantic Kernel "maintenance mode" — overstatement that should be softened. Microsoft's official position (from Semantic Kernel blog by Shawn Henry, Principal GPM) is: "Both projects will remain supported but most investment is now focused on Microsoft Agent Framework." Semantic Kernel will be supported "for at least one year after Microsoft Agent Framework leaves Preview and is Generally Available." Microsoft explicitly says they will "continue to address critical bugs, security issues" and take "some existing features to GA" but

the "majority of new features will be built for Microsoft Agent Framework." Microsoft Agent Framework **Microsoft has not used the term "maintenance mode" officially** — this characterization appears in third-party analysis. The paper should say "de facto wind-down" or "investment shift" rather than "maintenance mode." [vendor documentation]

Google ADK v1.27.2: Confirmed accurate. Released March 17, 2026 on PyPI (`google_adk-1.27.2.tar.gz`). Apache 2.0 licensed. PyPI ADK 2.0 Alpha with graph-based workflows has also been announced. Google Available in Python, Java, TypeScript, and Go. [vendor documentation]

OpenAI Agents SDK v0.12.3: Slightly outdated. The current version as of March 19, 2026 is **v0.12.5**, not v0.12.3. PyPI Version 0.12.3 likely existed as an intermediate release. The v0.12.x series added sessions (Redis, SQLAlchemy, Dapr), realtime/voice agents, SIP protocol support, and gpt-5.4 model support for computer use. PyPI OpenAI MIT-licensed. [vendor documentation]

Topic 5: SGLang deployment claims are self-reported but corroborated

400,000+ GPUs: Self-reported by the SGLang project README, repeated across PyPI and third-party articles. No independent audit exists, but the figure is plausible given the organization evidence below. [vendor/project documentation]

Organization-by-organization validation: xAI has **strong** evidence — the PyTorch official blog states "xAI uses SGLang to serve its flagship model, Grok 3," and a presentation by Ying Sheng (xAI/LMSYS/UCLA) claims SGLang "serves the production and research workloads at xAI (100K+ GPUs)." Llmsystem AMD and NVIDIA both have **strong** evidence through dedicated documentation, container releases, and co-hosted meetups. LinkedIn and Cursor are **listed in the README only** — no independent blog posts or case studies were found confirming their specific deployments. [vendor documentation + practitioner evidence]

RadixAttention 5–6× throughput: Confirmed by peer-reviewed evidence. The SGLang paper (arXiv 2312.07104, accepted at NeurIPS) reports **"up to 6.4× higher throughput"** compared to state-of-the-art inference systems. The LMSYS blog reports "up to 5× faster." These are **peak/best-case figures for workloads with shared prefixes** (few-shot, multi-turn, agent chains). Real-world improvements are more modest — a RunPod benchmark found **10–20% improvement** over vLLM in realistic multi-turn scenarios. The "5–6×" claim is accurate as stated but should note it is a peak figure. [peer-reviewed]

SGLang v0.5.9: Confirmed. Released February 23–24, 2026 on both PyPI and GitHub. github This is the latest stable release as of March 20, 2026. [vendor documentation]

Topic 6: All four GPU prices confirmed exact

Every price point validates precisely against ComputePrices.com (last updated March 19, 2026) and corroborating sources:

Claim	Verified Price	Status
H100 SXM $2.69/hr RunPod	**$2.69/hr** (Secure Cloud)	✅ Exact match
H100 SXM $1.91/hr median Vast.ai	**$1.55–$2.27/hr** range (midpoint $1.91)	✅ Consistent
RTX 4090 $0.34/hr RunPod	**$0.34/hr**	✅ Exact match
H200 $3.59/hr RunPod	**$3.59/hr**	✅ Exact match

64–75% collapse from peak: Confirmed by multiple independent sources. Peak H100 rental prices in 2023 were **$7–$12+/hr**. The Introl blog (January 2026) calculates a **64% decline** from $8/hr to $2.85–$3.50/hr. Accio independently reports **"64–75% reduction from peak in Q4 2024."** Thunder Compute (March 2026) documents the drop from "$8/hour historical peaks to $1.90–$3.90." Causes include supply expansion, 300+ new GPU cloud providers entering the market, AWS's ~44% price cut in June 2025, the shift from training to inference workloads, and H200/Blackwell displacement pressure. [vendor documentation + industry analysis]

Note: RunPod Community Cloud offers even lower pricing (H100 SXM at ~$1.50/hr). computeprices The paper's cited prices appear to be Secure Cloud (dedicated) pricing. [vendor documentation]

Topic 7: All three claims confirmed

Open WebUI: 127,665 stars → currently ~127,842 stars. The paper's figure was accurate at time of writing; the repo continues gaining stars. GitHub The **290M+ downloads** figure matches Open WebUI's homepage claim ("290 million downloads and growing") Open WebUI and is corroborated by ByteByteGo citing "over 282 million downloads" in a slightly earlier article. ByteByteGo Downloads are distributed across GHCR container pulls and PyPI, not a single Docker Hub page. Current version is v0.8.10 (March 9, 2026). Releasealert PyPI [vendor documentation + practitioner evidence]

LibreChat acquired by ClickHouse: Confirmed, November 4, 2025. Announced via BusinessWire, ClickHouse blog, and multiple outlets. Founder Danny Avila and team joined ClickHouse. LibreChat remains **100% open-source under MIT.** ClickHouse The combined stack forms the "open-source Agentic Data Stack." ClickHouse ClickHouse's internal "DWAINE" agent

(powered by LibreChat + ClickHouse) handles ~70% of data warehouse queries for 200+ users, processing **33M LLM tokens daily**. Business Wire Current LibreChat version: v0.8.3-rc1, ~34,300 GitHub stars, 23M+ container pulls. [vendor documentation]

LobeHub rebranded from LobeChat: Confirmed, November 3, 2025. The GitHub repo was renamed from `lobehub/lobe-chat` to `lobehub/lobehub`. The official announcement states the rebrand reflects evolution from "conversations" to a "hub connecting users, diverse agents, vast knowledge, and external services." GitHub Part of the LobeHub 2.0 initiative including multi-agent collaboration, Agent Builder, and personal memory features. GitHub GitHub Currently at **~71,800 GitHub stars**. github [vendor documentation]

Topic 8: Ten developments that would strengthen the paper

1. Open models now reach 70–80% on SWE-bench Verified

This is the single biggest gap. As of February 2026, multiple open-weight models perform competitively with frontier closed models on software engineering tasks. **MiniMax M2.5 leads at 80.2%,** Vertu followed by Qwen 3.5 at **76.4%,** Vertu GLM-4.7 at **73.8%,** Vertu and Step-3.5-Flash at **74.4%** (with only 196B parameters). Vertu For context, Claude Opus 4.6 (Thinking) leads at 79.2% and GPT 5.4 at 77.2%. Vals AI The gap between open and closed models on coding tasks has **narrowed dramatically** since Llama 4's release. SWE-bench Pro (Scale Labs) introduces harder, contamination-resistant tasks where even frontier models drop to ~23%. Scale [peer-reviewed + vendor benchmarks]

2. OWASP published two critical agentic security frameworks

OWASP Top 10 for Agentic Applications (released December 2025, peer-reviewed by 100+ experts) OWASP identifies ten risk categories: Agent Goal Hijack, Tool Misuse, Identity & Privilege Abuse, Agentic Supply Chain Vulnerabilities, Unexpected Code Execution, Memory & Context Poisoning, Insecure Inter-Agent Communication, Cascading Failures, Human-Agent Trust Exploitation, and Rogue Agents. OWASP Referenced by Microsoft, NVIDIA, and AWS. OWASP A **separate OWASP MCP Top 10** (beta) covers MCP-specific risks including tool poisoning OWASP (84.2% success rate with auto-approval), DEV Community rug pulls, and supply chain compromise. A key finding: **an audit of 17 popular MCP servers found an average security score of 34/100, with 100% lacking permission declarations.** Over 30 CVEs were filed in 60 days. DEV Community MintMCP [community-vetted framework]

3. MCP security research has exploded beyond Invariant Labs

The paper should significantly expand its MCP security coverage. Key additions: **Palo Alto Networks Unit 42** (December 2025) demonstrated resource theft, conversation hijacking, and covert tool invocation via MCP sampling. Palo Alto Networks **JFrog discovered CVE-2025-6514 —**

a critical (CVSS 9.6) command injection in `mcp-remote` (eSentire) with 437k+ downloads. (arXiv) (Heyuan110) **Anthropic's own `mcp-server-git`** had 3 chained CVEs achieving full RCE. (Vulnerablemcp) Academic papers include "Model Context Protocol: Landscape, Security Threats, and Future Research Directions" (arXiv:2503.23278, October 2025) (arXiv) and the **MCPTox benchmark** (arXiv:2508.14925) — the first public MCP tool poisoning benchmark, finding **attack success rates above 60%** for GPT-4o-mini, o1-mini, DeepSeek-R1, and Phi-4. (arXiv) A dedicated vulnerability tracker exists at vulnerablemcp.info. [peer-reviewed + vendor security research]

4. GraphRAG performance is more nuanced than assumed

Two important peer-reviewed findings challenge a simple "GraphRAG is better" narrative. **"RAG vs. GraphRAG: A Systematic Evaluation"** (Han et al., MSU + Meta, arXiv:2502.11371) (arXiv) finds that (arxiv) **neither approach universally dominates** — GraphRAG excels at multi-hop reasoning (Medium) while standard RAG is better for simple factual retrieval. (arXiv) More critically, **GraphRAG-Bench (ICLR 2026)** reports that "recent studies report that GraphRAG frequently underperforms vanilla RAG on many real-world tasks." (GitHub) Quantitative evidence is mixed: Lettria/AWS found GraphRAG at **80% correct vs. 50.8% for standard RAG** on complex datasets, (AWS) but Cognilium documented **2.4× higher latency** as the trade-off for 1.5× better accuracy. (Cognilium) **Microsoft's LazyGraphRAG** outperformed all comparison conditions in 96/96 comparisons with statistical significance, (Microsoft) suggesting lighter-weight graph approaches may be the practical sweet spot. [peer-reviewed + vendor benchmarks]

5. Agentic RAG shows promise but is LLM-dependent

The key paper is **"Is Agentic RAG worth it?"** (arXiv:2601.07711, January 2026), comparing naive, enhanced, and agentic RAG across multiple LLMs. The critical finding: **agentic RAG performance is highly dependent on the underlying LLM quality** — weaker generators show disproportionate degradation. (arXiv) Contextual AI's RAG Agent achieved 71.2% on RAG-QA Arena, a **5.4% improvement** over the strongest baseline (Cohere + Claude-3.5-Sonnet at 66.8%). (Contextual AI) Context windows expanding to 1–2M tokens are also challenging fundamental RAG assumptions (AIMultiple) — some tasks are now better served by direct context processing than complex retrieval agents. [peer-reviewed + vendor benchmarks]

6. Post-Llama 4 model releases transformed the landscape

The paper needs a section covering major releases from April 2025 through March 2026. **Qwen 3/3.5** has **overtaken Llama in total downloads** and become the most-used base model for fine-tuning — Qwen 3.5 (397B) achieves 88.4% on GPQA Diamond and 76.4% on SWE-bench. (Vertu) **Kimi K2** (Moonshot AI) is "arguably the best open model by benchmark score" and leads on SWE-rebench. **GLM-4.7** achieves HumanEval 94.2% and AIME 2025 95.7%. (Vertu) **Mistral Large 3** is the first Mistral MoE (675B total / 41B active). (MashBlog) **Step-3.5-Flash** achieves AIME 97.3% at only 196B parameters. (Vertu) Chinese labs (DeepSeek, Qwen, Moonshot, Z.AI,

MiniMax, StepFun) now **dominate open-weight model development,** Understanding AI with downloads shifting from US-dominant to China-dominant during summer 2025. Red Hat [vendor documentation + practitioner evidence]

7. vLLM has evolved dramatically to v0.17

Current vLLM version is **v0.17.1** (March 11, 2026) PyPI — the paper should update from whatever version it references. Major developments include **vLLM-Omni** (November 2025) for multimodal serving (diffusion, image/video, audio/TTS), GitHub **vLLM Semantic Router v0.1 "Iris"** (January 2026) with production-ready semantic routing and HaluGate hallucination detection, vLLM and the V1 architecture rewrite deprecating the V0 engine. Hardware support has expanded to NVIDIA, AMD ROCm, Intel XPU, TPU, Huawei Ascend, IBM Spyre, and ARM CPUs. **Red Hat acquired Neural Magic** (vLLM's commercial backer), integrating vLLM into OpenShift AI. [vendor documentation]

8. EU AI Act full applicability arrives August 2, 2026

The paper should flag that **full EU AI Act enforcement begins August 2, 2026** European Commission — less than five months away. Most provisions including high-risk AI classification, transparency rules, regulatory sandboxes in every member state, and full enforcement powers activate on this date. Programming Helper Open-source AI systems receive explicit exemptions from many requirements unless classified as high-risk, unacceptable risk, or requiring transparency. Linuxfoundation Currently **no open-source models are classified as posing systemic risk**. Hugging Face A proposed Digital Omnibus (November 2025) would extend some high-risk deadlines to December 2027. nadcab.com [official EU documentation]

9. Agent framework ecosystem matured with six production-grade options

Beyond the frameworks the paper already covers, the landscape now includes **Anthropic Claude Agent SDK** (v0.1.48), **CrewAI** (v1.10.1 with native MCP + A2A support, 44,600+ stars), Let's Data Science and newer entrants like **Pydantic AI**, **Smolagents** (Hugging Face), **Strands Agents** Langfuse (AWS), **Agno**, and **Mastra**. All major frameworks now support MCP integration. AIMultiple **Google's A2A (Agent-to-Agent) protocol** enables cross-framework agent communication GuruSup — this is a significant architectural development worth covering. [vendor documentation]

10. Context windows challenge fundamental RAG assumptions

Models now support **1–2 million token context windows**, forcing a fundamental question the paper should address: when does direct context processing outperform complex retrieval architectures? This affects the paper's RAG recommendations and should be discussed as a counterpoint to the agentic RAG trend. [practitioner/community evidence]

Corrections summary

Claim	Status	Action needed
Llama 4 Maverick ELO 1417 controversy	✅ Confirmed + stronger	Add LeCun "fudged" confirmation (Jan 2026)
LangGraph "changed to proprietary"	❌ **Incorrect**	Correct to three-tier: library MIT, API ELv2, platform proprietary
MCP 97M downloads, 10K+ servers	✅ Confirmed	No change needed
AAIF / Linux Foundation Dec 2025	✅ Confirmed	No change needed
MCP Invariant Labs demos	✅ Confirmed	Expand with OWASP, CVEs, academic research
MCP SSE → Streamable HTTP	✅ Confirmed	Spec version 2025-03-26 deprecated SSE
Microsoft Agent Framework RC1 Feb 2026	⚠️ Mostly correct	Change "RC1" to "RC"; GA target is end Q1 2026
AutoGen + Semantic Kernel in maintenance	⚠️ Overstatement	Soften to "de facto wind-down" — not official Microsoft terminology
Google ADK v1.27.2	✅ Confirmed	No change needed
OpenAI Agents SDK v0.12.3	⚠️ Slightly outdated	Update to v0.12.5 (March 19, 2026)
SGLang 400K+ GPUs	⚠️ Self-reported	Note it is project-claimed; xAI/AMD/NVIDIA individually confirmed
RadixAttention 5–6×	✅ Confirmed (peer-reviewed)	Note this is peak/best-case for shared-prefix workloads
SGLang v0.5.9	✅ Confirmed	No change needed
H100 $2.69/hr RunPod	✅ Exact match	No change needed
H100 $1.91/hr Vast.ai median	✅ Consistent	No change needed
RTX 4090 $0.34/hr	✅ Exact match	No change needed

Claim	Status	Action needed
H200 $3.59/hr	✅ Exact match	No change needed
64–75% price collapse	✅ Confirmed	No change needed
Open WebUI 127,665 stars / 290M downloads	✅ Confirmed (~127,842 now)	No change needed
LibreChat acquired by ClickHouse	✅ Confirmed (Nov 4, 2025)	No change needed
LobeHub rebranded from LobeChat	✅ Confirmed (Nov 3, 2025)	No change needed
~87 validated registry servers	⚠️ Plausible, unverifiable	Note this is approximate and fluctuating
Scout VRAM at Q4_K_M	Reference: **~65 GB** per Ollama	Ensure paper states this clearly — common misconception that 17B active = small VRAM